

Byzantine Generals problem

Amirreza Taghizadeh

May 3, 2023

In this presentation, we aim to discuss a problem in distributed systems known as "**Byzantine generals problem**" proposed by **Leslie Lamport**.


Outline

- 1 Brief history of Lamport's works
- 2 Origination of Byzantine Generals problem
- 3 A review of the Byzantine problem
- 4 Signed messages
- 5 Practical Byzantine-fault tolerance
- 6 Topics Discussed

Outline

- 1 Brief history of Lamport's works
- 2 Origination of Byzantine Generals problem
- 3 A review of the Byzantine problem
- 4 Signed messages
- 5 Practical Byzantine-fault tolerance
- 6 Topics Discussed


Brief history of Lamport's works



- **Inventing \LaTeX :** a group of macros, which can make the life of \TeX users a lot easier!!
- **One way authentication** in Whitfield Diffie's "*New directions in cryptography*" (1976)

brief description of **One-way Function**

- A one way function **f** is a function that is **easy to compute** but whose **inverse is difficult to compute**:
- or to say f is one-way function iff (adapted from¹):
 - ① for all value v , finding data object d s.t. $\phi(d) = v$ is **computationally infeasible**.
 - ② f is not **one-to-one** or proving it otherwise is **computationally infeasible**.²



¹Lamport, "Constructing Digital Signatures from a One Way Function"

²i.e. to say: \forall data object $d' \neq d : \phi(d') \neq \phi(d)$

brief description of **One-way Authentication**

Definition (Digital signature)

A digital signature created by **sender P** for **document m** is a data item $\sigma_P(m)$ that is when received together with **m**, one can determine (e.g. in a court of law) that **P** generated document **m**.

Hence A tool for determining **validity** of something sent.¹

¹Lamport, L. (1979) " *Constructing Digital Signatures from a One Way Function*"

brief description of **One-way Authentication**

Definition (One way authentication)

It must be **easy for anyone** to recognize the signature as **authentic** but **impossible** for anyone other than the signer to produce it!¹

¹Diffie, W. (1976)" *New Directions in Cryptography*"

A practical Example of a **One-way function** in **one-way authentication**

Login Problem

- User **A** enters Password PW and computer store it as $f(PW)$
- where $f(PW)$ is a one-way function of 10 million instructions
- and its inverse has 10^{30} more instructions (or computations), which practically makes it **noninvertible**
- for example, finding square root of x_0 given in $f(x) = x^2$ is much harder than computing x^2 at x_0 .

brief description of **One-way Authentication** *Cont'd*

- However, determining exactly what the one-way function should be is originally solved by **Lamport**
which further lead to the publication of the paper: "*Constructing Digital Signatures from a One Way Function*"
- But how this solution relates to the ecosystem of **public keys** is out of the scope of the presentation and discussed in the paper: "*New Directions in Cryptography*" by Whitfield Diffie (1976)

Brief history of Lamport's works

- **Inventing L^AT_EX:** a group of macros, which can make the life of T_EX users a lot easier!!
- **One way authentication** in Whitfield Diffie's "*New directions in cryptography*" (1976)
- **bakery algorithm** an algorithm to ensure mutual exclusion in concurrent processes
- that is to ensure a data structure is modified by at most one process at a time
and no process is reading a data structure while it is being written by other processes.
- **Paxos algorithm:** an algorithm used in distributed systems for reaching consensus, used in distributed storage systems

¹Silberschatz, "Database System Concepts", Ch. 19, P. 965

Brief history of Lamport's works *cont'd*

- Time, clocks and ordering of events in a distributed system:
 - ▶ in a distributed system, sometimes it is **impossible** to say an event has happened before something else.
 - ▶ hence, relation "*happened before*" is a partial ordering relation.
 - ▶ **Partial ordering? sounds familiar...**

Definition (Partial ordering)

Partial ordering relation is an ordering relation in which not all members of the set need to be comparable!

- ▶ He proposed in that paper a partial ordering of "*happened before*" and gave a **distributed algorithm** for **extending it to a total ordering of events**

But where is the "Byzantine generals" problem in the list?

Outline

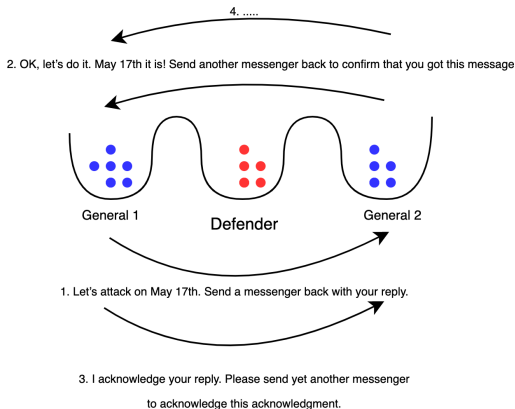
- 1 Brief history of Lamport's works
- 2 Origination of Bezyntine Generals problem**
- 3 A review of the Byzantine problem
- 4 Signed messages
- 5 Practical Byzantine-fault tolerance
- 6 Topics Discussed

Brief history of the origination of the problem

- **Naming:** "Because it was posed as a cute problem about philosophers seated around a table, **Dijkstra's dining philosopher's problem** received much more attention than it deserves. (For example, it has probably received more attention in the theory community than the readers/writers problem, which illustrates the same principles and has much more practical importance.) . . . The popularity of the dining philosophers problem taught me that **the best way to attract attention to a problem is to present it in terms of a story.**" -Lamport
- **Motivation:** Two General's problem. Byzantine General is a rather more general form of this problem.

A brief Overview of Two general's problem

Stating the problem through visualization!



Note: All Acknowledgement messages were sent by assuming the defender not capturing any messengers

Outline

- 1 Brief history of Lamport's works
- 2 Origination of Byzantine Generals problem
- 3 A review of the Byzantine problem
- 4 Signed messages
- 5 Practical Byzantine-fault tolerance
- 6 Topics Discussed

A review of the Byzantine problem

- There is an **enemy city** and a group of **General i** , each deciding to reach **an agreed upon plan** (which is the exact definition of **consensus**) to whether *Attack* or *Retreat*
- and each general i is equipped with a messaging method for sending value $v(i) (\in \{\text{Attack}, \text{Retreat}\})$ for communicating with each other
- AND there are a bunch of **traitorous generals** sending **conflicting** messages, aim to prevent **loyal generals to reach a plan**
- In fact, they send **arbitrary messages** to other generals.
- In order for them to reach consensus, two conditions must be satisfied:
 - 1 Every loyal general must obtain the same information
 $v(1), v(2) \dots, v(n)$
 - 2 The value sent by a loyal general should be used by all loyal generals

A review of the Byzantine problem

A review of the Byzantine problem *cont.*

How should the generals send their messages?

Let's examine how **a single general i should send the message $v(i)$** that is formally defined by: *why??* (which is made by grouping generals into two groups, namely **commander and lieutenant generals**)

Definition (Byzantine Generals Problem)

A **commanding general** must send an order to his $n - 1$ **lieutenant generals** s.t.:

IC1. All loyal lieutenants obey the same order.

IC2. If the commanding general is loyal, then every loyal lieutenant obeys the order he sends.

Note that $IC2 \implies IC1$

* Also note that, E.g. when General n sends $v(n)$ lieutenant $n - 1$ retrieves a message from $n - 2$ generals and then apply a function $Majority(v(1), v(2), \dots, v(n - 2))$ and adds $v(n)$ to the list V_1

A review of the Byzantine problem *cont.*

Lamport gave a recursive algorithm based on **majority function** for the mentioned problem in case of having **oral messages** whose content is solely managed by the sender.

Unfortunately, the algorithm there are two problems concerning this algorithm:

- 1 **It is expensive** $\rightarrow O(n!)$
- 2 It only works only, in case of having m traitors, for $n \geq 3m + 1$ generals

In order to deal with problem 2, he proposed an algorithm based on **unforgeable messages**.

Outline


- 1 Brief history of Lamport's works
- 2 Origination of Byzantine Generals problem
- 3 A review of the Byzantine problem
- 4 Signed messages**
- 5 Practical Byzantine-fault tolerance
- 6 Topics Discussed

Signed messages

What was wrong with the oral messages?

Answer: because **traitors lie**, they **alter the contents** of the messages they receive and send.

Let's make messages that can not be lied (forged) or any alterations could be detected



Formal & abstract Assumptions regarding Messages!

- A1. Every Message that is sent is delivered correctly.
 - A2. The receiver of a message knows who sent it.
 - A3. The absence of a message can be detected.
- } → **Oral messages**

A4. (a) A loyal general's signature cannot be forged (**or changed!!**) and any alteration of the contents of his signed messages can be detected

(b) Anyone can verify the authenticity of a general's signature



Signed messages

BUT, in case of assuming A4, What is the purpose of a traitor?????

Signed messages algorithm

Note that we know **the number of m** traitors when running the Alg.

- ① Commander sends a **message** having a value v and a signature (which is a sequence of IDs) $\rightarrow v : 0$
- ② each lieutenant i receives the message of length k ,
adds the v to a V_i set, adds his ID to the message and sends it to those child lieutenants not having received this message before
- ③ when lieutenant i receives no more messages, the lieutenant i applies the **a Choice function** to V_i in order to retrieve an order.

What is that **Choice function**?

Choice function

The Choice function could be any **aggregate function** (such as median, average, etc) BUT, it needs to have two essential properties:

- 1 if Set V_i consists of single value v Then, $Choice(V_i) = v$
- 2 $Choice(\emptyset) = RETREAT$

Formal statement of Signed messages $SM(m)$

Initially $V_i = \emptyset$

(1) The commander signs and sends message $v : 0$ to all lieutenants

(2) For each i :

(A) If Lieutenant i receives a message of the form $v : 0$ from the commander and he hasn't received any order, then:

(i) he lets $V_i = v$

(ii) he sends the message $v : 0 : i$ to every other lieutenant.


(B) If Lieutenant i receives a message of the form $v : 0 : j_1 \cdots : j_k$ **and** $v \notin V_i$ then:

(i) he adds v to V_i ;

(ii) if $k < m$, then he sends the message $v : 0 : j_1 \cdots : j_k : i$ to every lieutenant other than j_1, \dots, j_k

(3) For each i : When Lieutenant i will receive no more messages, he obeys the order $Choice(V_i)$

The impossible case revisited



- The proposed algorithm **increased fault tolerance**, but it is still **expensive** $\rightarrow O(n!)$
- In what follows, we give a description of a more **practical and efficient algorithm** which can tolerate Byzantine-faults (i.e. arbitrary messages and faults in nodes) in asynchronous systems, proposed by *Castro M. & Liskov B.*
- but first let us define what is meant by **synchronous and asynchronous systems**.

Outline

- 1 Brief history of Lamport's works
- 2 Origination of Byzantine Generals problem
- 3 A review of the Byzantine problem
- 4 Signed messages
- 5 Practical Byzantine-fault tolerance**
- 6 Topics Discussed

Definition (Synchronous distributed system)

A distributed system comprising of processes is synchronous iff it has the following properties¹

- 1 There is a real-time physical clock which provides synchronization among multiple processes
- 2 A known upper bound ϵ on processing delays
- 3 A known upper bound δ on message transmission delays.

^a "Reliable and Secure Distributed Programming" A book by Rodrigues L. & Cachin

Asynchronous systems → Lamport's Time clocks revisited

Definition (Asynchronous system¹)

In such systems, there is no time assumptions about processes. That is, we do not have **physical clock**. Instead **Time** is defined with respect to communication and measured by a *Logical clock*.

Here is Lamport's Algorithm for finding something has "happened before" something else.

¹"*Reliable and Secure Distributed Programming*" A book by Rodrigues L. & Cachin

A brief description of Lamport's Algorithm for Time clocks

- 1 Each process p keeps an integer called *logical clock* l_p initially 0.
- 2 Whenever an event occurs at Node p , the logical clock l_p is incremented by one unit.
- 3 When a process **sends a message**, it adds a timestamp to the message with the value of its logical clock at the moment the message is sent. The timestamp of an event e is denoted by $t(e)$.
- 4 When a Node p receives a message m with timestamp t_m , Node p increments its logical clock in the following way: $l_p := \max\{l_p, t_m\} + 1$

A brief description of Lamport's Algorithm for Time clocks *cont.*

- now that we have a mechanism for time in asynchronous world, let us define the *happened before* relation

Definition (Happened before relation $e_1 \rightarrow e_2$)


For event e_1 and e_2 we say " e_1 has happened before e_2 " ($e_1 \rightarrow e_2$) if e_1 and e_2 have the following conditions:

- 1 e_1 and e_2 occurred at the **same process p** and e_1 occurred before e_2 (i.e. $t(e_2) = t(e_1) + 1$)
- 2 e_1 corresponds to the transmission of a message **m** at process p and e_2 to the reception of m at process q .
- 3 \rightarrow be transitive.

¹"Reliable and Secure Distributed Programming" A book by Rodrigues L. & Cachin

A brief description of Lamport's Algorithm for Time clocks *cont.*

The above definition implies that if $e_1 \rightarrow e_2 \implies t(e_1) < t(e_2)$



Practical Byzantine Fault Tolerance (PBFT)

This Algorithm describes the implementation of a **Byzantine-fault-tolerant distributed file system**.

Assumptions:

- 1 We have a **asynchronous distributed system** \implies network may have: **Delays, failure to deliver messages; duplicate them, delivery out of order**
- 2 For Byzantine assumption (nodes behaving arbitrarily): Independent node failures \implies each node runs service code **independently**, have different root password and a different admin.
- 3 Use cryptographic techniques to prevent spoofing and to detect corrupted messages \rightarrow messages have *public key signatures*, message authentication codes.
- 4 A strong adversary coordinates faulty nodes, delay communication and delay correct nodes
- 5 The adversary cannot violate cryptographic techniques used by nodes.

Practical Byzantine Fault Tolerance (PBFT)

Properties:

- ① The algorithm used to implement any **replicated service** with *state* and *some operations*, including but not limited to *read and writes on replicas*
- ② There are some **faulty** and **non-faulty clients**(**requesting to invoke operations**) and **replicas** (we have f faulty replicas and $n = 3f + 1$ number of all replicas))
- ③ The Algorithm both provides **safety** and **liveness**
 - ① **Safety (Accuracy)**: replicated service satisfies linearizability.
 - ② **Liveness (Completeness)**: Clients eventually receive replies to their requests, provided at most $\lfloor \frac{n-1}{3} \rfloor$ are faulty and $delay(t)$ does not grow faster than t indefinitely.

- ④ Algorithm does not rely on **synchrony** to provide **safety**
- ⑤ However, **It must** rely on synchrony to provide **liveness** → **FLP's impossibility**
- ⑥ The algorithm uses a 3-phase commit protocol: Pre-prepare, prepare, commit

PBFT Algorithm

- There is a primary replica in each request of a client c .
- Replicas move through a succession of configurations called **views**.
- in a view, one replica is the primary and others are backups.

PBFT Alg. Cont.

A very brief description of the Algorithm:

- A client sends a request to invoke a service operation to the primary replica
- The client's request message is of the form $\langle REQUEST, o, t, c \rangle_{\sigma_c}$
- The primary multicasts the request to the backups
- Replicas execute the request and send a reply to the client
- The client waits for $f + 1$ replies from different replicas with the same result; this the result of the operation.
- The replica's reply is of the form $\langle REQUEST, \nu, t, c, i, r \rangle_{\sigma_i}$

PBFT Alg. Cont.

- The client waits for $f + 1$ replies with valid signatures from different replicas, and with the same t and r , before accepting the result r
- If client doesn't receive replies soon, it broadcasts the request to all replicas.
- If the request has already been processed, the replicas simply re-send the reply;
- otherwise, if the replica is not the primary, it relays the request to the primary.
- If the primary doesn't multicast the request to the group, it will eventually be suspected to be faulty by enough replicas to cause a view change.
- **The algorithm is proved to be Only 3% slower than non-replicated implementation of NFS**

Outline

- 1 Brief history of Lamport's works
- 2 Origination of Bezyntine Generals problem
- 3 A review of the Byzantine problem
- 4 Signed messages
- 5 Practical Byzantine-fault tolerance
- 6 Topics Discussed**

Topics Discussed

- Appreciation of Lamport's Other Works: \LaTeX , Paxos, Bakery algorithm, One-way authentication, **Time clocks in Asynchronous Distributed systems** etc.
- What is the Byzantine generals problem and how does it relate to the problem of solving **consensus**?
- Using Oral messaging to solve Byzantine-fault problem + **its weaknesses** $\rightarrow n \geq 3m + 1$
- What is the structure of **signed and unforgeable messages** and How does it solve Oral Messaging?
- Defining time with **logical clocks** in a system **without a global physical clock** (Asynchronous Systems)
- Is there a more efficient algorithm for Solving Byzantine failure in asynchronous System? \rightarrow Practical Byzantine-fault Tolerant Algorithm

References:



L. Lamport, R. Shostak, and M. Pease

The Byzantine Generals Problem

ACM Transactions on Programming Languages and Systems (1982).



M. Fishcher, N. Lynch, and M. Paterson

Impossibility of Distributed Consensus with One Faulty Process

Journal of the Association for Computing Machinery, Vol. 32, No. 2, (1985).



C. Cachin, R. Guerraoui and L. Rodrigues

Introduction to Reliable and Secure Distributed Programming, 2nd Edition

Springer (2011), p. 44-48.



A. Silberschatz, H. Korth, and S. Sudarshan

Database System Concepts, 7th Edition

McGraw-Hill Education (2020), p. 965



L. Lamport

Time, Clocks, and the Ordering of Events in a Distributed System

Journal of the Association for Computing Machinery, Vol. 21, No. 7, (1978)

References:



M. Castro, B. Liskov

Practical Byzantine Fault Tolerance

Third Symposium on Operating Systems Design and Implementation, New Orleans, (1999)



W. Diffie and M. Hellman

New Directions in Cryptography

IEEE Transactions on Information Theory, Vol. IT-22, No. 6 (1976)



L. Lamport

Constructing Digital Signatures from a One Way Function

SRI International (1979)



F. A. Parand

Failure and Consensus

(.ppt) file used as teaching material in Distributed Systems class by Prof. Parand, Allameh Tabatabai University

Any Questions?