

# Byzantine Generals problem

Amirreza Taghizadeh

April 27, 2023

In this presentation, we aim to discuss a problem in distributed systems known as "**Byzantine generals problem**" proposed by **Leslie Lamport**.


# Outline

- 1 Brief history of Lamport's works
- 2 Origination of Byzantine Generals problem
- 3 A review of the Byzantine problem
- 4 Signed messages
- 5 Solving consensus in asynchronous and synchronous systems

# Outline

- 1 Brief history of Lamport's works
- 2 Origination of Byzantine Generals problem
- 3 A review of the Byzantine problem
- 4 Signed messages
- 5 Solving consensus in asynchronous and synchronous systems

# Brief history of Lamport's works



- **Inventing  $\text{\LaTeX}$ :** a group of macros, which can make the life of  $\text{\TeX}$ users a lot easier!!
- **One way authentication** in Whitfield Diffie's "*New directions in cryptography*" (1976)

## brief description of **One-way Function**

- A one way function  **$f$**  is a function that is **easy to compute** but whose **inverse is difficult to compute**:

---

<sup>1</sup>Lamport, " *Constructing Digital Signatures from a One Way Function*"

<sup>2</sup>i.e. to say:  $\forall$  data object  $d' \neq d : \phi(d') \neq \phi(d)$

## brief description of **One-way Function**

- A one way function  **$f$**  is a function that is **easy to compute** but whose **inverse is difficult to compute**:
- or to say  $f$  is one-way function iff (adapted from<sup>1</sup>):

---

<sup>1</sup>Lamport, " *Constructing Digital Signatures from a One Way Function*"

<sup>2</sup>i.e. to say:  $\forall$  data object  $d' \neq d : \phi(d') \neq \phi(d)$

## brief description of **One-way Function**

- A one way function  **$f$**  is a function that is **easy to compute** but whose **inverse is difficult to compute**:
- or to say  $f$  is one-way function iff (adapted from<sup>1</sup>):
  - ① for all value  $v$ , finding data object  $d$  s.t.  $\phi(d) = v$  is **computationally infeasible**.

---

<sup>1</sup>Lamport, "Constructing Digital Signatures from a One Way Function"

<sup>2</sup>i.e. to say:  $\forall$  data object  $d' \neq d : \phi(d') \neq \phi(d)$



## brief description of **One-way Function**

- A one way function  **$f$**  is a function that is **easy to compute** but whose **inverse is difficult to compute**:
- or to say  $f$  is one-way function iff (adapted from<sup>1</sup>):
  - ① for all value  $v$ , finding data object  $d$  s.t.  $\phi(d) = v$  is **computationally infeasible**.
  - ②  $f$  is not **one-to-one** or proving it otherwise is **computationally infeasible**.<sup>2</sup>


---

<sup>1</sup>Lamport, "Constructing Digital Signatures from a One Way Function"

<sup>2</sup>i.e. to say:  $\forall$  data object  $d' \neq d : \phi(d') \neq \phi(d)$

# brief description of **One-way Function**

- A one way function  **$f$**  is a function that is **easy to compute** but whose **inverse is difficult to compute**:
- or to say  $f$  is one-way function iff (adapted from<sup>1</sup>):
  - ① for all value  $v$ , finding data object  $d$  s.t.  $\phi(d) = v$  is **computationally infeasible**.
  - ②  $f$  is not **one-to-one** or proving it otherwise is **computationally infeasible**.<sup>2</sup>



<sup>1</sup>Lamport, "Constructing Digital Signatures from a One Way Function"


<sup>2</sup>i.e. to say:  $\forall$  data object  $d' \neq d : \phi(d') \neq \phi(d)$

# brief description of **One-way Authentication**

## Definition (Digital signature)

A digital signature created by **sender P** for **document m** is a data item  $\sigma_P(m)$  that is when received together with **m**, one can determine (e.g. in a court of law) that **P** generated document **m**.

---

<sup>1</sup>Lamport, L. (1979) " *Constructing Digital Signatures from a One Way Function*" 


# brief description of **One-way Authentication**

## Definition (Digital signature)

A digital signature created by **sender P** for **document m** is a data item  $\sigma_P(m)$  that is when received together with **m**, one can determine (e.g. in a court of law) that **P** generated document **m**.

Hence A tool for determining **validity** of something sent.<sup>1</sup>

---

<sup>1</sup>Lamport, L. (1979) " *Constructing Digital Signatures from a One Way Function*" 

# brief description of **One-way Authentication**

## Definition (One way authentication)

It must be **easy for anyone** to recognize the signature as **authentic** but **impossible** for anyone other than the signer to produce it!<sup>1</sup>

---

<sup>1</sup>Diffie, W. (1976)" *New Directions in Cryptography*"

# A practical Example of a **One-way function** in **one-way authentication**

## Login Problem

- User **A** enters Password  $PW$  and computer store it as  $f(PW)$

# A practical Example of a **One-way function** in **one-way authentication**

## Login Problem

- User **A** enters Password  $PW$  and computer store it as  $f(PW)$
- where  $f(PW)$  is a one-way function of 10 million instructions

# A practical Example of a **One-way function** in **one-way authentication**

## Login Problem

- User **A** enters Password  $PW$  and computer store it as  $f(PW)$
- where  $f(PW)$  is a one-way function of 10 million instructions
- and its inverse has  $10^{30}$  more instructions (or computations), which practically makes it **noninvertible**



# A practical Example of a **One-way function** in **one-way authentication**

## Login Problem

- User **A** enters Password  $PW$  and computer store it as  $f(PW)$
- where  $f(PW)$  is a one-way function of 10 million instructions
- and its inverse has  $10^{30}$  more instructions (or computations), which practically makes it **noninvertible**
- for example, finding square root of  $x_0$  given in  $f(x) = x^2$  is much harder than computing  $x^2$  at  $x_0$ .

## brief description of **One-way Authentication** *Cont'd*

- However, determining exactly what the one-way function should be is originally solved by **Lamport**  
which further lead to the publication of the paper: "*Constructing Digital Signatures from a One Way Function*"

## brief description of **One-way Authentication** *Cont'd*

- But how this solution relates to the ecosystem of **public keys** is out of the scope of the presentation and discussed in the paper: *"New Directions in Cryptography"* by Whitfield Diffie (1976)

# Brief history of Lamport's works

- **Inventing  $\text{\LaTeX}$ :** a group of macros, which can make the life of  $\text{\TeX}$  users a lot easier!!
- **One way authentication** in Whitfield Diffie's "*New directions in cryptography*" (1976)
- **bakery algorithm** an algorithm to ensure mutual exclusion in concurrent processes

---

<sup>1</sup>Silberschatz, "Database System Concepts", Ch. 19, P. 965

# Brief history of Lamport's works

- **Inventing  $\text{\LaTeX}$ :** a group of macros, which can make the life of  $\text{\TeX}$  users a lot easier!!
- **One way authentication** in Whitfield Diffie's "*New directions in cryptography*" (1976)
- **bakery algorithm** an algorithm to ensure mutual exclusion in concurrent processes
- that is to ensure a data structure is modified by at most one process at a time  
and no process is reading a data structure while it is being written by other processes.

---

<sup>1</sup>Silberschatz, "Database System Concepts", Ch. 19, P. 965

# Brief history of Lamport's works

- **Inventing  $\text{\LaTeX}$ :** a group of macros, which can make the life of  $\text{\TeX}$  users a lot easier!!
- **One way authentication** in Whitfield Diffie's "*New directions in cryptography*" (1976)
- **bakery algorithm** an algorithm to ensure mutual exclusion in concurrent processes
- that is to ensure a data structure is modified by at most one process at a time  
and no process is reading a data structure while it is being written by other processes.
- **Paxos algorithm:** an algorithm used in distributed systems for reaching consensus, used in distributed storage systems

---

<sup>1</sup>Silberschatz, "Database System Concepts", Ch. 19, P. 965

## Brief history of Lamport's works *cont'd*

- Time, clocks and ordering of events in a distributed system:

# Brief history of Lamport's works *cont'd*

- Time, clocks and ordering of events in a distributed system:
  - ▶ in a distributed system, sometimes it is **impossible** to say an event has happened before something else.



# Brief history of Lamport's works *cont'd*

- Time, clocks and ordering of events in a distributed system:
  - ▶ in a distributed system, sometimes it is **impossible** to say an event has happened before something else.
  - ▶ hence, relation "*happened before*" is a partial ordering relation.

## Brief history of Lamport's works *cont'd*

- Time, clocks and ordering of events in a distributed system:
  - ▶ in a distributed system, sometimes it is **impossible** to say an event has happened before something else.
  - ▶ hence, relation "*happened before*" is a partial ordering relation.
  - ▶ **Partial ordering? sounds familiar...**

## Brief history of Lamport's works *cont'd*

- Time, clocks and ordering of events in a distributed system:
  - ▶ in a distributed system, sometimes it is **impossible** to say an event has happened before something else.
  - ▶ hence, relation "*happened before*" is a partial ordering relation.
  - ▶ **Partial ordering? sounds familiar...**

### Definition (Partial ordering)

Partial ordering relation is an ordering relation in which not all members of the set need to be comparable!

## Brief history of Lamport's works *cont'd*

- Time, clocks and ordering of events in a distributed system:
  - ▶ in a distributed system, sometimes it is **impossible** to say an event has happened before something else.
  - ▶ hence, relation "*happened before*" is a partial ordering relation.
  - ▶ He proposed in that paper a partial ordering of "*happened before*" and gave a **distributed algorithm** for **extending it to a total ordering of events**

## Brief history of Lamport's works *cont'd*

- Time, clocks and ordering of events in a distributed system:
  - ▶ in a distributed system, sometimes it is **impossible** to say an event has happened before something else.
  - ▶ hence, relation "*happened before*" is a partial ordering relation.
  - ▶ He proposed in that paper a partial ordering of "*happened before*" and gave a **distributed algorithm** for **extending it to a total ordering of events**

**But where is the "Byzantine generals" problem in the list?**

# Outline

- 1 Brief history of Lamport's works
- 2 Origination of Bezyntine Generals problem
- 3 A review of the Byzantine problem
- 4 Signed messages
- 5 Solving consensus in asynchronous and synchronous systems

# Brief history of the origination of the problem

- **Naming:** "As the Edsger W. Dijkstra's "Dining philosophers (a classic problem discussed in Operating Systems classes) involves a story, I decided to **include a story with the problem** related to Digital signatures with a recursive solution algorithm" -Lamport

# Brief history of the origination of the problem

- **Naming:** "As the Edsger W. Dijkstra's "Dining philosophers (a classic problem discussed in Operating Systems classes) involves a story, I decided to **include a story with the problem** related to Digital signatures with a recursive solution algorithm" -Lamport
- **Motivation:** Two General's problem.

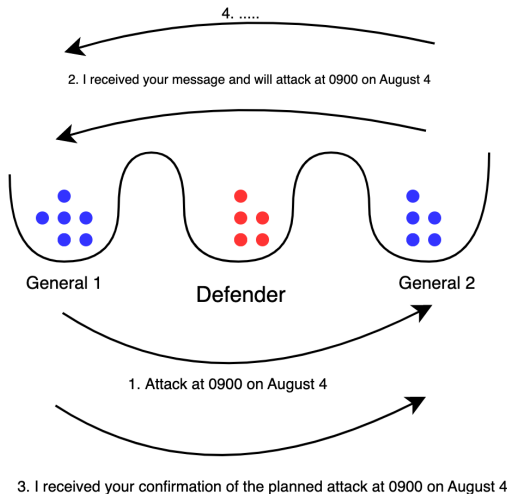


# Brief history of the origination of the problem

- **Naming:** "As the Edsger W. Dijkstra's "Dining philosophers (a classic problem discussed in Operating Systems classes) involves a story, I decided to **include a story with the problem** related to Digital signatures with a recursive solution algorithm" -Lamport
- **Motivation:** Two General's problem. in fact, byzantine generals are a more general form of this problem.

# A brief Overview of Two general's problem

## Stating the problem through visualization!



# Outline

- 1 Brief history of Lamport's works
- 2 Origination of Byzantine Generals problem
- 3 A review of the Byzantine problem**
- 4 Signed messages
- 5 Solving consensus in asynchronous and synchronous systems

# A review of the Byzantine problem

- There is an **enemy city** and a group of **General  $i$** , each deciding to reach **an agreed upon plant** (which is the exact definition of **consensus**)

# A review of the Byzantine problem

- There is an **enemy city** and a group of **General  $i$** , each deciding to reach **an agreed upon plant** (which is the exact definition of **consensus**)
- and each general  $i$  is equipped with a messaging method for sending value  $v(i)$

# A review of the Byzantine problem

- There is an **enemy city** and a group of **General  $i$** , each deciding to reach **an agreed upon plan** (which is the exact definition of **consensus**)
- and each general  $i$  is equipped with a messaging method for sending value  $v(i)$
- AND there are a bunch of **traitorous generals** sending **conflicting** messages, aim to prevent **loyal generals to reach a plan**

# A review of the Byzantine problem

- There is an **enemy city** and a group of **General  $i$** , each deciding to reach **an agreed upon plan** (which is the exact definition of **consensus**)
- and each general  $i$  is equipped with a messaging method for sending value  $v(i)$
- AND there are a bunch of **traitorous generals** sending **conflicting** messages, aim to prevent **loyal generals to reach a plan**
- In order for them to reach consensus, two conditions must be satisfied:

# A review of the Byzantine problem

- There is an **enemy city** and a group of **General  $i$** , each deciding to reach **an agreed upon plan** (which is the exact definition of **consensus**)
- and each general  $i$  is equipped with a messaging method for sending value  $v(i)$
- AND there are a bunch of **traitorous generals** sending **conflicting** messages, aim to prevent **loyal generals to reach a plan**
- In order for them to reach consensus, two conditions must be satisfied:
  - 1 Every loyal general must obtain the same information  $v(1), v(2) \dots, v(n)$
  - 2 The value sent by a loyal general should be used by all loyal generals



# A review of the Byzantine problem *cont*

How should the generals send their messages?

## A review of the Byzantine problem *cont*

How should the generals send their messages?

Let's examine how **a single general  $i$  should send the message  $v(i)$**  that is formally defined by:

## A review of the Byzantine problem *cont*

How should the generals send their messages?

Let's examine how **a single general  $i$  should send the message  $v(i)$**  that is formally defined by: (which is made by grouping generals into two groups, namely **commander and lieutenant generals**)

### Definition (Byzantine Generals Problem)

A **commanding general** must send an order to his  $n - 1$  **lieutenant generals** s.t.:

IC1. All loyal lieutenants obey the same order.

IC2. If the commanding general is loyal, then every loyal lieutenant obeys the order he sends.

## A review of the Byzantine problem *cont*

How should the generals send their messages?

Let's examine how **a single general  $i$  should send the message  $v(i)$**  that is formally defined by: (which is made by grouping generals into two groups, namely **commander and lieutenant generals**)

### Definition (Byzantine Generals Problem)

A **commanding general** must send an order to his  $n - 1$  **lieutenant generals** s.t.:

IC1. All loyal lieutenants obey the same order.

IC2. If the commanding general is loyal, then every loyal lieutenant obeys the order he sends.

Note that  $IC2 \implies IC1$

## A review of the Byzantine problem *cont*

Lamport gave a recursive algorithm based on **majority function** for the mentioned problem in case of having **oral messages** whose content is solely managed by the sender.

## A review of the Byzantine problem *cont*

Lamport gave a recursive algorithm based on **majority function** for the mentioned problem in case of having **oral messages** whose content is solely managed by the sender.

Unfortunately, the algorithm works only, in case of having  $m$  traitors, for  $n \geq 3m + 1$  generals

## A review of the Byzantine problem *cont*

Lamport gave a recursive algorithm based on **majority function** for the mentioned problem in case of having **oral messages** whose content is solely managed by the sender.

Unfortunately, the algorithm works only, in case of having  $m$  traitors, for  $n \geq 3m + 1$  generals

This is where he proposed an algorithm based on  
**unforgeable messages.**

# Outline

- 1 Brief history of Lamport's works
- 2 Origination of Byzantine Generals problem
- 3 A review of the Byzantine problem
- 4 Signed messages**
- 5 Solving consensus in asynchronous and synchronous systems



# Signed messages

What was wrong with the oral messages??

# Signed messages


What was wrong with the oral messages??

Answer: because **traitors lie**, they **alter the contents** of the messages they receive and send

# Signed messages

What was wrong with the oral messages??

Answer: because **traitors lie**, they **alter the contents** of the messages they receive and send



# Signed messages

What was wrong with the oral messages??

Answer: because **traitors lie**, they **alter the contents** of the messages they receive and send

Let's make messages that can not be lied (forged) or any alterations could be detected

# Signed messages algorithm

Note that we know **the number of  $m$**  traitors when running the Alg.

# Signed messages algorithm

- 1 Commander sends a message having a value  $v$  and a signature (which is a sequence of IDs)  $\rightarrow \mathbf{v : 0}$

# Signed messages algorithm

- ① Commander sends a message having a value  $v$  and a signature (which is a sequence of IDs)  $\rightarrow v : 0$
- ② each lieutenant  $i$  receives the message of length  $k$ ,  
**adds the  $v$  to a  $V_i$  set, adds his ID to the message and sends it to those child lieutenants not having received this message before**

# Signed messages algorithm

- 1 Commander sends a message having a value  $v$  and a signature (which is a sequence of IDs)  $\rightarrow v : 0$
- 2 each lieutenant  $i$  receives the message of length  $k$ ,  
**adds the  $v$  to a  $V_i$  set, adds his ID to the message and sends it to those child lieutenants not having received this message before**
- 3 when lieutenant  $i$  receives no more messages, the lieutenant  $i$  applies the **a Choice function** to  $V_i$  in order to retrieve an order.



# Signed messages algorithm

- 1 Commander sends a message having a value  $v$  and a signature (which is a sequence of IDs)  $\rightarrow v : 0$
- 2 each lieutenant  $i$  receives the message of length  $k$ ,  
**adds the  $v$  to a  $V_i$  set, adds his ID to the message and sends it to those child lieutenants not having received this message before**
- 3 when lieutenant  $i$  receives no more messages, the lieutenant  $i$  applies the **a Choice function** to  $V_i$  in order to retrieve an order.

what is that **Choice function**?

# Choice function

The Choice function could be any **aggregate function** (such as median, average, etc) BUT, it needs to have two essential properties:

- 1 if Set  $V_i$  consists of single value  $v$  Then,  $Choice(V_i) = v$

# Choice function

The Choice function could be any **aggregate function** (such as median, average, etc) BUT, it needs to have two essential properties:

- 1 if Set  $V_i$  consists of single value  $v$  Then,  $Choice(V_i) = v$
- 2  $Choice(\emptyset) = RETREAT$

# Formal statement of Signed messages

Initially  $V_i = \emptyset$

(1) The commander signs and sends message  $v : 0$  to all lieutenants

(2) For each  $i$ :

(A) If Lieutenant  $i$  receives a message of the form  $v : 0$  from the commander and he hasn't received any order, then:

(i) he lets  $V_i = v$

(ii) he sends the message  $v : 0 : i$  to every other lieutenant.


(B) If Lieutenant  $i$  receives a message of the form  $v : 0 : j_1 \cdots : j_k$  **and**  $v \notin V_i$  then:

(i) he adds  $v$  to  $V_i$ ;

(ii) if  $k < m$ , then he sends the message  $v : 0 : j_1 \cdots : j_k : i$  to every lieutenant other than  $j_1, \dots, j_k$

(3) For each  $i$ : When Lieutenant  $i$  will receive no more messages, he obeys the order  $Choice(V_i)$

## The impossible case revisited



# Outline

- 1 Brief history of Lamport's works
- 2 Origination of Byzantine Generals problem
- 3 A review of the Byzantine problem
- 4 Signed messages
- 5 Solving consensus in asynchronous and synchronous systems

Through the end of the presentation, we will be talking about **Consensus in asynchronous** systems and an important Article in this regard, by Fischer-Lynch-Paterson (FLP) Which proves the impossibility of **reaching consensus in asynchronous system** in the presence of any failure. However, in the paper

# Asynchronous systems → Lamport's Time clocks revisited

## Definition (Asynchronous system<sup>1</sup>)

In such systems, there is no time assumptions about processes. That is, we do not have **physical clock**. Instead **Time** is defined with respect to communication and measured by a *Logical clock*.

---

<sup>1</sup>"*Reliable and Secure Distributed Programming*" A book by Rodrigues L. & Cachin C.



# Asynchronous systems → Lamport's Time clocks revisited

## Definition (Asynchronous system<sup>1</sup>)

In such systems, there is no time assumptions about processes. That is, we do not have **physical clock**. Instead **Time** is defined with respect to communication and measured by a *Logical clock*.

Here is Lamport's Algorithm for finding something has "happened before" something else.

---

C. <sup>1</sup>"*Reliable and Secure Distributed Programming*" A book by Rodrigues L. & Cachin

# Lamport's Algorithm for Time clocks

- 1 Each process  $p$  keeps an integer called *logical clock*  $I_p$  initially 0.

---

C. <sup>1</sup>"*Reliable and Secure Distributed Programming*" A book by Rodrigues L. & Cachin

# Lamport's Algorithm for Time clocks

- 1 Each process  $p$  keeps an integer called *logical clock*  $I_p$  initially 0.
- 2 Whenever an event occurs at Node  $p$ , the logical clock  $I_p$  is incremented by one unit.

---

C. <sup>1</sup>"*Reliable and Secure Distributed Programming*" A book by Rodrigues L. & Cachin

# Lamport's Algorithm for Time clocks

- ① Each process  $p$  keeps an integer called *logical clock*  $I_p$  initially 0.
- ② Whenever an event occurs at Node  $p$ , the logical clock  $I_p$  is incremented by one unit.
- ③ When a process **sends a message**, it adds a timestamp to the message with the value of its logical clock at the moment the message is sent. The timestamp of an event  $e$  is denoted by  $t(e)$ .

---

C. <sup>1</sup>"*Reliable and Secure Distributed Programming*" A book by Rodrigues L. & Cachin

# Lamport's Algorithm for Time clocks

- 1 Each process  $p$  keeps an integer called *logical clock*  $I_p$  initially 0.
- 2 Whenver an event occurs at Node  $p$ , the logical clock  $I_p$  is incremented by one unit.
- 3 When a process **sends a message**, it adds a timestamp to the message with the value of its logical clock at the moment the message is sent. The timestamp of an event  $e$  is denoted by  $t(e)$ .
- 4 When a Node  $p$  receives a message  $m$  with timestamp  $t_m$ , Node  $p$  increments its logical clock in the following way:  $I_p := \max\{I_p, t_m\} + 1$

---

C. <sup>1</sup>"Reliable and Secure Distributed Programming" A book by Rodrigues L. & Cachin