

Definition (Synchronous distributed system)

A distributed system comprising of processes is synchronous iff it has the following properties¹

- 1 There is a real-time physical clock which provides synchronization among multiple processes
- 2 A known upper bound ϵ on processing delays
- 3 A known upper bound δ on message transmission delays.

^a"Reliable and Secure Distributed Programming" A book by Rodrigues L. & Cachin

Definition (Synchronous distributed system)

A distributed system comprising of processes is synchronous iff it has the following properties¹

- 1 There is a real-time physical clock which provides synchronization among multiple processes
- 2 A known upper bound ϵ on processing delays
- 3 A known upper bound δ on message transmission delays.

^a"Reliable and Secure Distributed Programming" A book by Rodrigues L. & Cachin

Asynchronous systems → Lamport's Time clocks revisited

Definition (Asynchronous system²)

In such systems, there is no time assumptions about processes. That is, we do not have **physical clock**. Instead **Time** is defined with respect to communication and measured by a *Logical clock*.

Asynchronous systems → Lamport's Time clocks revisited

Definition (Asynchronous system²)

In such systems, there is no time assumptions about processes. That is, we do not have **physical clock**. Instead **Time** is defined with respect to communication and measured by a *Logical clock*.

¹"Reliable and Secure Distributed Programming" A book by Rodrigues L. & Cachin

¹"Reliable and Secure Distributed Programming" A book by Rodrigues L. & Cachin

Asynchronous systems → Lamport's Time clocks revisited

Definition (Asynchronous system²)

In such systems, there is no time assumptions about processes. That is, we do not have **physical clock**. Instead **Time** is defined with respect to communication and measured by a *Logical clock*.

Here is Lamport's Algorithm for finding something has "happened before" something else.

Asynchronous systems → Lamport's Time clocks revisited

Definition (Asynchronous system²)

In such systems, there is no time assumptions about processes. That is, we do not have **physical clock**. Instead **Time** is defined with respect to communication and measured by a *Logical clock*.

Here is Lamport's Algorithm for finding something has "happened before" something else.

¹"Reliable and Secure Distributed Programming" A book by Rodrigues L. & Cachin

¹"Reliable and Secure Distributed Programming" A book by Rodrigues L. & Cachin

A brief description of Lamport's Algorithm for Time clocks

- 1 Each process p keeps an ineger called *logical clock* l_p initially 0.

A brief description of Lamport's Algorithm for Time clocks

- 1 Each process p keeps an ineger called *logical clock* l_p initially 0.
-

A brief description of Lamport's Algorithm for Time clocks

- 1 Each process p keeps an ineger called *logical clock* I_p initially 0.
- 2 Whenever an event occurs at Node p , the logical clock I_p is incremented by one unit.

A brief description of Lamport's Algorithm for Time clocks

- 1 Each process p keeps an ineger called *logical clock* I_p initially 0.
 - 2 Whenever an event occurs at Node p , the logical clock I_p is incremented by one unit.
-

A brief description of Lamport's Algorithm for Time clocks

- 1 Each process p keeps an ineger called *logical clock* I_p initially 0.
- 2 Whenver an event occurs at Node p , the logical clock I_p is incremented by one unit.
- 3 When a process **sends a message**, it adds a timestamp to the message with the value of its logical clock at the moment the message is sent. The timestamp of an event e is denoted by $t(e)$.

A brief description of Lamport's Algorithm for Time clocks

- 1 Each process p keeps an ineger called *logical clock* I_p initially 0.
 - 2 Whenver an event occurs at Node p , the logical clock I_p is incremented by one unit.
 - 3 When a process **sends a message**, it adds a timestamp to the message with the value of its logical clock at the moment the message is sent. The timestamp of an event e is denoted by $t(e)$.
-

A brief description of Lamport's Algorithm for Time clocks

- 1 Each process p keeps an ineger called *logical clock* l_p initially 0.
- 2 Whenver an event occurs at Node p , the logical clock l_p is incremented by one unit.
- 3 When a process **sends a message**, it adds a timestamp to the message with the value of its logical clock at the moment the message is sent. The timestamp of an event e is denoted by $t(e)$.
- 4 When a Node p receives a message m with timestamp t_m , Node p increments its logical clock in the following way: $l_p := \max\{l_p, t_m\} + 1$

A brief description of Lamport's Algorithm for Time clocks

- 1 Each process p keeps an ineger called *logical clock* l_p initially 0.
 - 2 Whenver an event occurs at Node p , the logical clock l_p is incremented by one unit.
 - 3 When a process **sends a message**, it adds a timestamp to the message with the value of its logical clock at the moment the message is sent. The timestamp of an event e is denoted by $t(e)$.
 - 4 When a Node p receives a message m with timestamp t_m , Node p increments its logical clock in the following way: $l_p := \max\{l_p, t_m\} + 1$
-

A brief description of Lamport's Algorithm for Time clocks *cont.*

- now that we have a mechanism for time in asynchronous world, let us define the *happened before* relation

A brief description of Lamport's Algorithm for Time clocks *cont.*

- now that we have a mechanism for time in asynchronous world, let us define the *happened before* relation

¹"Reliable and Secure Distributed Programming" A book by Rodrigues L. & Cachin

A brief description of Lamport's Algorithm for Time clocks *cont.*

- now that we have a mechanism for time in asynchronous world, let us define the *happened before* relation

Definition (Happened before relation $e_1 \rightarrow e_2$)

For event e_1 and e_2 we say " e_1 has happened before e_2 " ($e_1 \rightarrow e_2$) if e_1 and e_2 have the following conditions:

- 1 e_1 and e_2 occurred at the **same process p** and e_1 occurred before e_2 (i.e. $t(e_2) = t(e_1) + 1$)
- 2 e_1 corresponds to the transmission of a message **m** at process p and e_2 to the reception of m at process q .
- 3 \rightarrow be transitive.

A brief description of Lamport's Algorithm for Time clocks *cont.*

- now that we have a mechanism for time in asynchronous world, let us define the *happened before* relation

Definition (Happened before relation $e_1 \rightarrow e_2$)

For event e_1 and e_2 we say " e_1 has happened before e_2 " ($e_1 \rightarrow e_2$) if e_1 and e_2 have the following conditions:

- 1 e_1 and e_2 occurred at the **same process p** and e_1 occurred before e_2 (i.e. $t(e_2) = t(e_1) + 1$)
- 2 e_1 corresponds to the transmission of a message **m** at process p and e_2 to the reception of m at process q .
- 3 \rightarrow be transitive.

¹"Reliable and Secure Distributed Programming" A book by Rodrigues L. & Cachin

-
- event is something such as sending or receiving message.
 - there exists some event e' s.t. $e_1 \rightarrow e'$ and $e' \rightarrow e_2$

¹"Reliable and Secure Distributed Programming" A book by Rodrigues L. & Cachin

A brief description of Lamport's Algorithm for Time clocks

cont.

The above definition implies that if $e_1 \rightarrow e_2 \implies t(e_1) < t(e_2)$

A brief description of Lamport's Algorithm for Time clocks

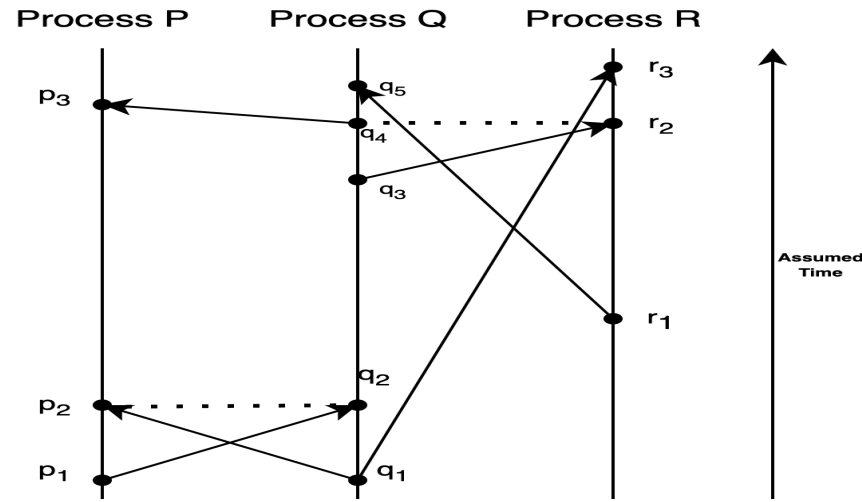
cont.

The above definition implies that if $e_1 \rightarrow e_2 \implies t(e_1) < t(e_2)$

A brief description of Lamport's Algorithm for Time clocks

cont.

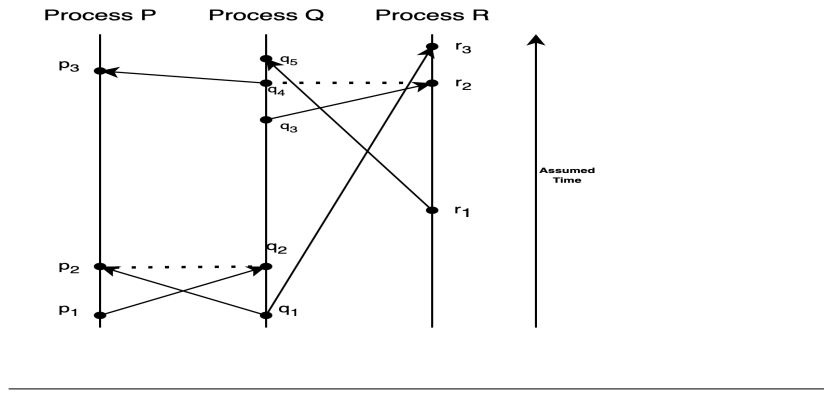
The above definition implies that if $e_1 \rightarrow e_2 \implies t(e_1) < t(e_2)$



A brief description of Lamport's Algorithm for Time clocks

cont.

The above definition implies that if $e_1 \rightarrow e_2 \implies t(e_1) < t(e_2)$



Practical Byzantine Fault Tolerance (PBFT)

This Algorithm describes the implementation of a
Byzantine-fault-tolerant distributed file system.

Assumptions:

- 1 We have a **asynchronous distributed system** \implies network may have: **Delays, failure to deliver messages; duplicate them, delivery out of order**

Practical Byzantine Fault Tolerance (PBFT)

This Algorithm describes the implementation of a
Byzantine-fault-tolerant distributed file system.

Assumptions:

- 1 We have a **asynchronous distributed system** \implies network may have: **Delays, failure to deliver messages; duplicate them, delivery out of order**

Practical Byzantine Fault Tolerance (PBFT)

This Algorithm describes the implementation of a **Byzantine-fault-tolerant distributed file system**.

Assumptions:

- 1 We have a **asynchronous distributed system** \implies network may have: **Delays, failure to deliver messages; duplicate them, delivery out of order**
- 2 For Byzantine assumption (nodes behaving arbitrarily): Independent node failures \implies each node runs service code **independently**, have different root password and a different admin.

Practical Byzantine Fault Tolerance (PBFT)

This Algorithm describes the implementation of a **Byzantine-fault-tolerant distributed file system**.

Assumptions:

- 1 We have a **asynchronous distributed system** \implies network may have: **Delays, failure to deliver messages; duplicate them, delivery out of order**
- 2 For Byzantine assumption (nodes behaving arbitrarily): Independent node failures \implies each node runs service code **independently**, have different root password and a different admin.

Practical Byzantine Fault Tolerance (PBFT)

This Algorithm describes the implementation of a **Byzantine-fault-tolerant distributed file system.**

Assumptions:

- 1 We have a **asynchronous distributed system** \implies network may have: **Delays, failure to deliver messages; duplicate them, delivery out of order**
- 2 For Byzantine assumption (nodes behaving arbitrarily): Independent node failures \implies each node runs service code **independently**, have different root password and a different admin.
- 3 Use cryptographic techniques to prevent spoofing and to detect corrupted messages \rightarrow messages have *public key signatures*, message authentication codes.

Practical Byzantine Fault Tolerance (PBFT)

This Algorithm describes the implementation of a **Byzantine-fault-tolerant distributed file system.**

Assumptions:

- 1 We have a **asynchronous distributed system** \implies network may have: **Delays, failure to deliver messages; duplicate them, delivery out of order**
- 2 For Byzantine assumption (nodes behaving arbitrarily): Independent node failures \implies each node runs service code **independently**, have different root password and a different admin.
- 3 Use cryptographic techniques to prevent spoofing and to detect corrupted messages \rightarrow messages have *public key signatures*, message authentication codes.

Practical Byzantine Fault Tolerance (PBFT)

This Algorithm describes the implementation of a **Byzantine-fault-tolerant distributed file system.**

Assumptions:

- 1 We have a **asynchronous distributed system** \implies network may have: **Delays, failure to deliver messages; duplicate them, delivery out of order**
- 2 For Byzantine assumption (nodes behaving arbitrarily): Independent node failures \implies each node runs service code **independently**, have different root password and a different admin.
- 3 Use cryptographic techniques to prevent spoofing and to detect corrupted messages \rightarrow messages have *public key signatures*, message authentication codes.
- 4 A strong adversary coordinates faulty nodes, delay communication and delay correct nodes
- 5 The adversary cannot violate cryptographic techniques used by nodes.

Practical Byzantine Fault Tolerance (PBFT)

This Algorithm describes the implementation of a **Byzantine-fault-tolerant distributed file system.**

Assumptions:

- 1 We have a **asynchronous distributed system** \implies network may have: **Delays, failure to deliver messages; duplicate them, delivery out of order**
- 2 For Byzantine assumption (nodes behaving arbitrarily): Independent node failures \implies each node runs service code **independently**, have different root password and a different admin.
- 3 Use cryptographic techniques to prevent spoofing and to detect corrupted messages \rightarrow messages have *public key signatures*, message authentication codes.
- 4 A strong adversary coordinates faulty nodes, delay communication and delay correct nodes
- 5 The adversary cannot violate cryptographic techniques used by nodes.

For example, the adversary cannot produce a valid signature of a non-faulty node, compute the information summarized by a digest from the digest, or find two messages with the same digest. The cryptographic techniques we use are thought to have these properties

Practical Byzantine Fault Tolerance (PBFT)

Properties:

- 1 The algorithm used to implement any **replicated service** with *state* and *some operations*, including but not limited to *read and writes on replicas*

Practical Byzantine Fault Tolerance (PBFT)

Properties:

- 1 The algorithm used to implement any **replicated service** with *state* and *some operations*, including but not limited to *read and writes on replicas*
-

Practical Byzantine Fault Tolerance (PBFT)

Properties:

- 1 The algorithm used to implement any **replicated service** with *state* and *some operations*, including but not limited to *read and writes on replicas*
- 2 There are some **faulty** and **non-faulty clients**(requesting to invoke **operations**) and **replicas** (we have f faulty replicas and $n = 3f + 1$ number of all replicas))

Practical Byzantine Fault Tolerance (PBFT)

Properties:

- 1 The algorithm used to implement any **replicated service** with *state* and *some operations*, including but not limited to *read and writes on replicas*
- 2 There are some **faulty** and **non-faulty clients**(requesting to invoke **operations**) and **replicas** (we have f faulty replicas and $n = 3f + 1$ number of all replicas))

Practical Byzantine Fault Tolerance (PBFT)

Properties:

- 1 The algorithm used to implement any **replicated service** with *state* and *some operations*, including but not limited to *read and writes on replicas*
- 2 There are some **faulty** and **non-faulty clients**(requesting to invoke **operations**) and **replicas** (we have f faulty replicas and $n = 3f + 1$ number of all replicas))

Practical Byzantine Fault Tolerance (PBFT)

Properties:

- 1 The algorithm used to implement any **replicated service** with *state* and *some operations*, including but not limited to *read and writes on replicas*
- 2 There are some **faulty** and **non-faulty clients**(requesting to invoke **operations**) and **replicas** (we have f faulty replicas and $n = 3f + 1$ number of all replicas))

Faulty clients are dealt with by **strong authentication**

Practical Byzantine Fault Tolerance (PBFT)

Properties:

- 1 The algorithm used to implement any **replicated service** with *state* and *some operations*, including but not limited to *read and writes on replicas*
- 2 There are some **faulty** and **non-faulty clients**(requesting to invoke **operations**) and **replicas** (we have f faulty replicas and $n = 3f + 1$ number of all replicas))
- 3 The Algorithm both provides **safety** and **liveness**

Practical Byzantine Fault Tolerance (PBFT)

Properties:

- 1 The algorithm used to implement any **replicated service** with *state* and *some operations*, including but not limited to *read and writes on replicas*
- 2 There are some **faulty** and **non-faulty clients**(requesting to invoke **operations**) and **replicas** (we have f faulty replicas and $n = 3f + 1$ number of all replicas))
- 3 The Algorithm both provides **safety** and **liveness**

Faulty clients are dealt with by **strong authentication**

Practical Byzantine Fault Tolerance (PBFT)

Properties:

- ① The algorithm used to implement any **replicated service** with *state* and *some operations*, including but not limited to *read and writes on replicas*
- ② There are some **faulty** and **non-faulty clients**(requesting to invoke **operations**) and **replicas** (we have f faulty replicas and $n = 3f + 1$ number of all replicas))
- ③ The Algorithm both provides **safety** and **liveness**
 - ① **Safety (Accuracy)**: replicated service satisfies linearizability.

Practical Byzantine Fault Tolerance (PBFT)

Properties:

- ① The algorithm used to implement any **replicated service** with *state* and *some operations*, including but not limited to *read and writes on replicas*
- ② There are some **faulty** and **non-faulty clients**(requesting to invoke **operations**) and **replicas** (we have f faulty replicas and $n = 3f + 1$ number of all replicas))
- ③ The Algorithm both provides **safety** and **liveness**
 - ① **Safety (Accuracy)**: replicated service satisfies linearizability.

Faulty clients are dealt with by **strong authentication**- It behaves like a centralized implementation that executes operations **atomically** one at a time.

Practical Byzantine Fault Tolerance (PBFT)

Properties:

- ① The algorithm used to implement any **replicated service** with *state* and *some operations*, including but not limited to *read and writes on replicas*
- ② There are some **faulty** and **non-faulty clients**(requesting to invoke **operations**) and **replicas** (we have f faulty replicas and $n = 3f + 1$ number of all replicas))
- ③ The Algorithm both provides **safety** and **liveness**
 - ① **Safety (Accuracy)**: replicated service satisfies linearizability.
 - ② **Liveness (Completeness)**: Clients eventually receive replies to their requests, provided at most $\lfloor \frac{n-1}{3} \rfloor$ are faulty and $delay(t)$ does not grow faster than t indefinitely.

Practical Byzantine Fault Tolerance (PBFT)

Properties:

- ① The algorithm used to implement any **replicated service** with *state* and *some operations*, including but not limited to *read and writes on replicas*
- ② There are some **faulty** and **non-faulty clients**(requesting to invoke **operations**) and **replicas** (we have f faulty replicas and $n = 3f + 1$ number of all replicas))
- ③ The Algorithm both provides **safety** and **liveness**
 - ① **Safety (Accuracy)**: replicated service satisfies linearizability.
 - ② **Liveness (Completeness)**: Clients eventually receive replies to their requests, provided at most $\lfloor \frac{n-1}{3} \rfloor$ are faulty and $delay(t)$ does not grow faster than t indefinitely.

Faulty clients are dealt with by **strong authentication**- It behaves like a centralized implementation that executes operations **atomically** one at a time.

- $delay(t)$ is the time between the moment when a message is sent for the first time and the moment when it is received by its destination (assuming the sender keeps retransmitting the message until it is received)

4 Algorithm does not rely on **synchrony** to provide **safety**

Algorithm does not rely on **synchrony** to provide **safety**

- ④ Algorithm does not rely on **synchrony** to provide **safety**
- ⑤ However, **It must** rely on synchrony to provide **liveness** → **FLP's impossibility**

- ④ Algorithm does not rely on **synchrony** to provide **safety**
 - ⑤ However, **It must** rely on synchrony to provide **liveness** → **FLP's impossibility**
-

- ④ Algorithm does not rely on **synchrony** to provide **safety**
- ⑤ However, **It must** rely on synchrony to provide **liveness** → **FLP's impossibility**
- ⑥ The algorithm uses a 3-phase commit protocol: Pre-prepare, prepare, commit

- ④ Algorithm does not rely on **synchrony** to provide **safety**
 - ⑤ However, **It must** rely on synchrony to provide **liveness** → **FLP's impossibility**
 - ⑥ The algorithm uses a 3-phase commit protocol: Pre-prepare, prepare, commit
-

- There is a primary replica in each request of a client c .

- There is a primary replica in each request of a client c .
-

- There is a primary replica in each request of a client c .
- Replicas move through a succession of configurations called **views**.

- There is a primary replica in each request of a client c .
 - Replicas move through a succession of configurations called **views**.
-

- There is a primary replica in each request of a client c .
- Replicas move through a succession of configurations called **views**.
- in a view, one replica is the primary and others are backups.

- There is a primary replica in each request of a client c .
 - Replicas move through a succession of configurations called **views**.
 - in a view, one replica is the primary and others are backups.
-

A very brief description of the Algorithm:

A very brief description of the Algorithm:

A very brief description of the Algorithm:

- A client sends a request to invoke a service operation to the primary replica

A very brief description of the Algorithm:

- A client sends a request to invoke a service operation to the primary replica

A very brief description of the Algorithm:

- A client sends a request to invoke a service operation to the primary replica
- The client's request message is of the form $\langle REQUEST, o, t, c \rangle_{\sigma_c}$

A very brief description of the Algorithm:

- A client sends a request to invoke a service operation to the primary replica
- The client's request message is of the form $\langle REQUEST, o, t, c \rangle_{\sigma_c}$

A very brief description of the Algorithm:

- A client sends a request to invoke a service operation to the primary replica
- The client's request message is of the form $\langle REQUEST, o, t, c \rangle_{\sigma_c}$

A very brief description of the Algorithm:

- A client sends a request to invoke a service operation to the primary replica
- The client's request message is of the form $\langle REQUEST, o, t, c \rangle_{\sigma_c}$

c is client, o is the operation needed, timestamp t is used to ensure **exactly once** semantics for the execution of client requests.

A very brief description of the Algorithm:

- A client sends a request to invoke a service operation to the primary replica
- The client's request message is of the form $\langle REQUEST, o, t, c \rangle_{\sigma_c}$
- The primary multicasts the request to the backups

A very brief description of the Algorithm:

- A client sends a request to invoke a service operation to the primary replica
- The client's request message is of the form $\langle REQUEST, o, t, c \rangle_{\sigma_c}$
- The primary multicasts the request to the backups

c is client, o is the operation needed, timestamp t is used to ensure **exactly once** semantics for the execution of client requests.

A very brief description of the Algorithm:

- A client sends a request to invoke a service operation to the primary replica
- The client's request message is of the form $\langle REQUEST, o, t, c \rangle_{\sigma_c}$
- The primary multicasts the request to the backups
- Replicas execute the request and send a reply to the client

A very brief description of the Algorithm:

- A client sends a request to invoke a service operation to the primary replica
- The client's request message is of the form $\langle REQUEST, o, t, c \rangle_{\sigma_c}$
- The primary multicasts the request to the backups
- Replicas execute the request and send a reply to the client

c is client, o is the operation needed, timestamp t is used to ensure **exactly once** semantics for the execution of client requests.

A very brief description of the Algorithm:

- A client sends a request to invoke a service operation to the primary replica
- The client's request message is of the form $\langle REQUEST, o, t, c \rangle_{\sigma_c}$
- The primary multicasts the request to the backups
- Replicas execute the request and send a reply to the client
- The client waits for $f + 1$ replies from different replicas with the same result; this the result of the operation.

A very brief description of the Algorithm:

- A client sends a request to invoke a service operation to the primary replica
- The client's request message is of the form $\langle REQUEST, o, t, c \rangle_{\sigma_c}$
- The primary multicasts the request to the backups
- Replicas execute the request and send a reply to the client
- The client waits for $f + 1$ replies from different replicas with the same result; this the result of the operation.

c is client, o is the operation needed, timestamp t is used to ensure **exactly once** semantics for the execution of client requests.

A very brief description of the Algorithm:

- A client sends a request to invoke a service operation to the primary replica
- The client's request message is of the form $\langle REQUEST, o, t, c \rangle_{\sigma_c}$
- The primary multicasts the request to the backups
- Replicas execute the request and send a reply to the client
- The client waits for $f + 1$ replies from different replicas with the same result; this the result of the operation.
- The replica's reply is of the form $\langle REQUEST, \nu, t, c, i, r \rangle_{\sigma_i}$

A very brief description of the Algorithm:

- A client sends a request to invoke a service operation to the primary replica
- The client's request message is of the form $\langle REQUEST, o, t, c \rangle_{\sigma_c}$
- The primary multicasts the request to the backups
- Replicas execute the request and send a reply to the client
- The client waits for $f + 1$ replies from different replicas with the same result; this the result of the operation.
- The replica's reply is of the form $\langle REQUEST, \nu, t, c, i, r \rangle_{\sigma_i}$

c is client, o is the operation needed, timestamp t is used to ensure **exactly once** semantics for the execution of client requests.

A very brief description of the Algorithm:

- A client sends a request to invoke a service operation to the primary replica
- The client's request message is of the form $\langle REQUEST, o, t, c \rangle_{\sigma_c}$
- The primary multicasts the request to the backups
- Replicas execute the request and send a reply to the client
- The client waits for $f + 1$ replies from different replicas with the same result; this the result of the operation.
- The replica's reply is of the form $\langle REQUEST, \nu, t, c, i, r \rangle_{\sigma_i}$

A very brief description of the Algorithm:

- A client sends a request to invoke a service operation to the primary replica
- The client's request message is of the form $\langle REQUEST, o, t, c \rangle_{\sigma_c}$
- The primary multicasts the request to the backups
- Replicas execute the request and send a reply to the client
- The client waits for $f + 1$ replies from different replicas with the same result; this the result of the operation.
- The replica's reply is of the form $\langle REQUEST, \nu, t, c, i, r \rangle_{\sigma_i}$

c is client, o is the operation needed, timestamp t is used to ensure **exactly once** semantics for the execution of client requests. v is the current view number, t is the timestamp of the corresponding request, i is the replica number, and r is the result of executing the requested operation

PBFT Alg. Cont.

- The client waits for $f + 1$ replies with valid signatures from different replicas, and with the same t and r , before accepting the result r

PBFT Alg. Cont.

- The client waits for $f + 1$ replies with valid signatures from different replicas, and with the same t and r , before accepting the result r
-

PBFT Alg. Cont.

- The client waits for $f + 1$ replies with valid signatures from different replicas, and with the same t and r , before accepting the result r
- If client doesn't receive replies soon, it broadcasts the request to all replicas.

PBFT Alg. Cont.

- The client waits for $f + 1$ replies with valid signatures from different replicas, and with the same t and r , before accepting the result r
 - If client doesn't receive replies soon, it broadcasts the request to all replicas.
-

PBFT Alg. Cont.

- The client waits for $f + 1$ replies with valid signatures from different replicas, and with the same t and r , before accepting the result r
- If client doesn't receive replies soon, it broadcasts the request to all replicas.
- If the request has already been processed, the replicas simply re-send the reply;

PBFT Alg. Cont.

- The client waits for $f + 1$ replies with valid signatures from different replicas, and with the same t and r , before accepting the result r
 - If client doesn't receive replies soon, it broadcasts the request to all replicas.
 - If the request has already been processed, the replicas simply re-send the reply;
-

- The client waits for $f + 1$ replies with valid signatures from different replicas, and with the same t and r , before accepting the result r
- If client doesn't receive replies soon, it broadcasts the request to all replicas.
- If the request has already been processed, the replicas simply re-send the reply;
- otherwise, if the replica is not the primary, it relays the request to the primary.

- The client waits for $f + 1$ replies with valid signatures from different replicas, and with the same t and r , before accepting the result r
 - If client doesn't receive replies soon, it broadcasts the request to all replicas.
 - If the request has already been processed, the replicas simply re-send the reply;
 - otherwise, if the replica is not the primary, it relays the request to the primary.
-

- The client waits for $f + 1$ replies with valid signatures from different replicas, and with the same t and r , before accepting the result r
- If client doesn't receive replies soon, it broadcasts the request to all replicas.
- If the request has already been processed, the replicas simply re-send the reply;
- otherwise, if the replica is not the primary, it relays the request to the primary.
- If the primary doesn't multicast the request to the group, it will eventually be suspected to be faulty by enough replicas to cause a view change.

- The client waits for $f + 1$ replies with valid signatures from different replicas, and with the same t and r , before accepting the result r
 - If client doesn't receive replies soon, it broadcasts the request to all replicas.
 - If the request has already been processed, the replicas simply re-send the reply;
 - otherwise, if the replica is not the primary, it relays the request to the primary.
 - If the primary doesn't multicast the request to the group, it will eventually be suspected to be faulty by enough replicas to cause a view change.
-

Thank you for your attention!

Thank you for your attention!
